



Kungliga Tekniska Högskolan
Valhallavägen 79
100 44 Stockholm

Parallel Computations for Large-Scale Problems

« Homework 2 »

February 16th - March 2nd



Authors

Rémi Domingues 920604-T239
Johan Wärnegård 920113-4914

Teacher

Michael Hanke

Scholar year 2014-2015

1 The broadcast operation is a one-to-all collective communication operation where one of the processes sends the same message to all other processes.

1.1 Design an algorithm for the broadcast operation using only point-to-point communications which requires only $O(\log P)$ communication steps

The following figure describes the steps of our algorithm. The showcase uses $P = 2^D$, but the algorithm is functional for any P.




FIGURE 1 – Communication broadcast algorithm in $O(\log P)$

This figure illustrates the following algorithm. For any p higher than 0, the process will receive its message from $p_{sender} = p_{receiver} - 2^D$ with D such that $2^D \leq p_{receiver} \leq 2^{D+1}$.

```

if  $p == 0$  then
     $D = 0$ 
     $msg = \text{"message"}$ 
else
     $D = \text{trunc}(\log_2(p))$ 
     $msg = \text{receive}(p - 2^D)$ 
     $++D$ 
end if

while  $p + 2^D \leq p$  do
     $\text{send}(msg, p + 2^D)$ 
     $++D$ 
end while

```

1.2 Do a (time-)performance analysis for your algorithm

The process involves no computation on any processor. We denote the communication time required to send the data set, consisting of n elements, between two processors t_{comm} . As usual $t_{comm} = t_{startup} + nt_{data}$. Each step in the algorithm takes time t_{comm} , in total $D = \log_2 P$ steps are required. Altogether this yields the total time as a function of P and n as :

$$T = (t_{startup} + nt_{data}) \cdot \log_2 P$$

1.3 How can the scatter operation be implemented using $O(\log P)$ communication steps?

Suppose we have a data set consisting of P elements known by a process p_0 . We want to scatter the data so that a process i holds one element of the data set, also named element i . We would do this using the broadcast algorithm previously defined, replacing some calls by the following :

```

if p == 0 then
    D = 0 size
    data = loadData()
else
    D = trunc(log2(p))                                ▷ cast from double to int
    data = receive(p - 2D)
    ++D
end if

while p + 2D ≤ p do
    array = split(data, 2)                               ▷ Split the dataset in two parts of equal size
    data = array[0]
    send(array[1], p + 2D)
    ++D
end while

```

After the first iteration, two processes will each hold one half of the data. After 2 iterations 4 processors hold one fourth each, and so forth until all P processors hold exactly one element. The number of iterations is :

$$D = \frac{\log(P)}{\log(2)} \Rightarrow T \propto \log_2(P)$$

- 2 Consider a matrix A distributed on a $P \times P$ process mesh. An algorithm has been given in the lecture for evaluating the matrix-vector product $y = Ax$. While x is column distributed, y is row distributed. In order to carry out a further multiplication Ay , the vector y must be transposed.

2.1 Design an algorithm for this transposition. You may use the results from problem 1.

In order to perform the transposition of the vector y , two steps are necessary (if we consider the current state as the one immediately after the internal multiplication of x_c by $A_{r,c}$ by $p_{r,c}$). We consider here that A is a square matrix (since we just performed the Ax operation and want to perform Ay) which is distributed on the $P \times P$ mesh grid. We also assume that $P = M$, with A a matrix of $M \times M$.




FIGURE 2 – Data distribution for each processor, with A a square matrix

Therefore, our transposition algorithm would first require an efficient information exchange between the processes of the same row to compute y_r (this should be achieved using recursive doubling with a complexity of $O(\log P)$ for the whole matrix), then an

axial symmetry communication to achieve the final transposition (we assume that each process knows its row r and column c).

```

    Let  $2^D < P < 2^{D+1}$ 
     $s = y_r$  ▷  $y_r$  is still incomplete and is only a part of the real  $y_r$ 
    if  $p \geq 2^D$  then
        send( $s$ , bitflip( $p$ ,  $p_{r,D}$ ))
    end if
    if  $p \leq P - 2^D$  then
        receive( $h$ , bitflip( $p$ ,  $p_{r,D}$ ))
         $s = s + h$ 
    end if
    if  $p \leq 2^D$  then
        for  $d = 0 : D-1$  do
            send( $s$ , bitflip( $p$ ,  $p_{r,D}$ ))
            receive( $h$ , bitflip( $p$ ,  $p_{r,D}$ ))
             $s = s + h$ 
        end for
    end if
    if  $p \leq P - 2^D$  then
        send( $s$ , bitflip( $p$ ,  $p_{r,D}$ ))
    end if
    if  $p \geq 2^D$  then
        receive( $s$ , bitflip( $p$ ,  $p_{r,D}$ ))
    end if
     $y_r = s$  ▷  $y_r$  has now been summed and is known by every process of the same row
    if  $r \neq c$  then ▷ No permutation required if the process is on the diagonal
        MPI_Sendrecv( $y_r$ ,  $p_{c,r}$ ,  $y_c$ ,  $p_{c,r}$ ) ▷  $y_r$  is sent and we retrieve  $y_r$ 
    end if

```

On the other hand, if we do not have a vector to sum and transpose but a square matrix of size $P \times P$ distributed on a $P \times P$ mesh grid, the algorithm is the following :

```

    if  $r \neq c$  then ▷ No permutation required if the process is on the diagonal
        MPI_Sendrecv( $y_r$ ,  $p_{c,r}$ ,  $y_c$ ,  $p_{c,r}$ ) ▷  $y_r$  is sent and we retrieve  $y_r$ 
    end if

```

2.2 Make a performance analysis

The complexity of the first algorithm (summing the parts of the vector \vec{y} then applying a transposition to prepare the system for the multiplication yA) is $O(\log(P) + 2)$.

The complexity of the second algorithm for a square matrix is $O(1)$.

3 Parallel implementation of the Jacobi iteration

See the following pages.